

# GRAPHAF: A FLOW-BASED AUTOREGRESSIVE MODEL FOR MOLECULAR GRAPH GENERATION

By: Chence Shi, Minkai Xu, Zhaocheng Zhu, Weinan Zhang, Ming Zhang, Jian Tang

Presented by: Daniel Ahn and Xuan Lin





# Introduction


- Molecular design with desired properties
  - Applications to drug discovery and materials science
  - Used to help find drugs for Covid-19 vaccine although recommendations were ultimately not used
- Large number of potential molecules due to number of atoms, types of bonds, orientations, etc.
  - Search space estimated to be as large as  $10^{33}$
- Amenable to machine learning approaches due to generative models and representation of chemical structure as graphs

## Related Works: Variational Autoencoder

- Autoencoders with a continuous latent space
- Force the latent space to be continuous by projecting input into properties of distributions
- Generates outputs by sampling from learned distributions
- Minimizes ELBO loss



what can happen without regularisation



what we want to obtain with regularisation



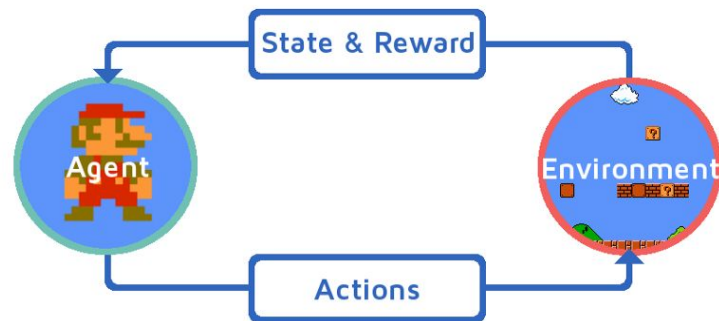
## Related Works: Generative Adversarial Network

- Generator and discriminator compete in a zero sum game
  - Generator must generate outputs that fool the discriminator
  - Discriminator must correctly distinguish between generated and real outputs
- No log likelihood maximization
  - Uses discriminator as a proxy for distance between learned and target distributions



## Related Works: Graph Convolutional Policy Network

- Problem reformulated in terms of states and actions with associated rewards
  - State: current graph structure
  - Action: new additions to graph
  - Reward: property of resulting graph
- Learn a correlation with graph state and next recommended addition by accruing reward for high performing final graphs





## Related Works: Autoregressive Models

- Predicted value on time step  $t$  is a function of values on timesteps  $T < t$

$$x_t = c + \sum_i^{t-1} \phi_i x_i + \epsilon$$

- Useful for modelling sequential data and processes



## Background: Normalizing Flow

- An invertible mapping

$$f : \mathcal{E} \rightarrow \mathcal{Z}$$

- where  $\mathcal{E}$  is the latent distribution and  $\mathcal{Z}$  is the data distribution
- Data distribution given by

$$p_Z(z) = p_{\mathcal{E}}(f_{\theta}^{-1}(z)) \left| \det \frac{\partial f_{\theta}^{-1}(z)}{\partial z} \right|.$$



## Background: Autoregressive Flow

- Latent distribution modeled by a gaussian over the previous time steps

$$p(z_d|z_{1:d-1}) = \mathcal{N}(z_d|\mu_d, (\alpha_d)^2), \text{ where } \mu_d = g_\mu(z_{1:d-1}; \theta), \alpha_d = g_\alpha(z_{1:d-1}; \theta).$$

- $g_\mu$  and  $g_\alpha$  are neural networks parameterized by  $\Theta$
- Data distribution calculated with the affine transformation

$$f_\theta(\epsilon_d) = z_d = \mu_d + \alpha_d \cdot \epsilon_d; \quad f_\theta^{-1}(z_d) = \epsilon_d = \frac{z_d - \mu_d}{\alpha_d}.$$





## Quiz

T/F: For autoregressive conditional probabilities parameterized by Gaussians, the mean and standard deviation of the latent distribution are only a function of the current sampled output

- True
- **False**



## Background: Graph Representation Learning

- Graph with  $n$  nodes,  $d$  types of nodes, and  $b$  types of edges is represented by  $G = (A, X)$ 
  - $A$  is the adjacency matrix  $A \in \{0, 1\}^{n \times n \times b}$ 
    - $A_{ijk} = 1$  if there is an edge of type  $k$  between nodes  $i$  and  $j$
  - $X$  is the node feature matrix  $X \in \{0, 1\}^{n \times d}$ 
    - $X_i = 1$  if the node has attribute  $i$
- A relational graph convolutional network (R-GCN) is used to learn node embeddings
  - Embedding at layer  $l$  is given by
$$H^l = \text{Agg} \left( \text{ReLU} \left( \{ \tilde{D}_i^{-\frac{1}{2}} \tilde{E}_i \tilde{D}_i^{-\frac{1}{2}} H^{l-1} W_i^l \} | i \in (1, \dots, b) \right) \right)$$
  - where  $\tilde{E} = A[:, :, i]$  with a self loop and  $\tilde{D}$  = the degree matrix of  $\tilde{E}$



## Methods: GraphAF

- Graph generation is modeled by a sequential process
  - Given graph  $G_i$ , a new node is added to the subgraph,  $p(X_i, G_i)$
  - After the nodes is added, an edge is added to connect the node to the rest of the graph,  $p(A_{ij}, G_i, X_i, A_{i,1:j-1})$ 
    - To allow for no edges to be connected, add a new edge type to indicate no connection
  - Repeat until either
    - Graph reaches maximum number of nodes
    - Newly added node cannot be connected to rest of the graph
- The autoregressive flow training process is parallelizable
  - Training is twice as fast on GraphAF compared to GCPN



## Methods: Dequantization

- Normalizing flow requires continuous inputs
  - A and X are discrete values so they must be transformed into continuous values
  - Use dequantization to transform  $G = (A, X)$  to  $z = (z^A, z^X)$

$$z_i^X = X_i + u, u \sim U[0, 1)^d; z_{ij}^A = A_{ij} + u, u \sim U[0, 1)^{b+1}$$

## Methods: Model Latent Distribution

- Apply autoregressive flow to the dequantized graph representation

- Latent distribution modeled by a gaussian

$$p(z_i^X | G_i) = \mathcal{N}(\mu_i^X, (\alpha_i^X)^2),$$

$$\text{where } \mu_i^X = g_{\mu^X}(G_i), \alpha_i^X = g_{\alpha^X}(G_i),$$

$$p(z_{ij}^A | G_i, X_i, A_{i,1:j-1}) = \mathcal{N}(\mu_{ij}^A, (\alpha_{ij}^A)^2), j \in \{1, 2, \dots, i-1\},$$

$$\text{where } \mu_{ij}^A = g_{\mu^A}(G_i, X_i, A_{i,1:j-1}), \alpha_{ij}^A = g_{\alpha^A}(G_i, X_i, A_{i,1:j-1}).$$

- In order to calculate the parameters  $\mu$  and  $\alpha$ , we must learn the neural networks  $g_\mu$  and  $g_\alpha$

- Use a R-GCN to get node embeddings from  $G_i$  and feed embeddings to MLPs

$$\text{R-GCN: } H_i^L = \text{R-GCN}(G_i), \tilde{h}_i = \text{sum}(H_i^L);$$

$$\text{Node-MLPs: } g_{\mu^X} = m_{\mu^X}(\tilde{h}_i), g_{\alpha^X} = m_{\alpha^X}(\tilde{h}_i);$$

$$\text{Edge-MLPs: } g_{\mu^A} = m_{\mu^A}(\tilde{h}_i, H_{i,i}^L, H_{i,j}^L), g_{\alpha^A} = m_{\alpha^A}(\tilde{h}_i, H_{i,i}^L, H_{i,j}^L)$$



## Methods: Sample and Predict

- Sample  $\epsilon_i$  and  $\epsilon_{ij}$  from the base Gaussian distribution
  - Apply affine transformation to recover the graph representation of the sample

$$z_i^X = \epsilon_i \odot \alpha_i^X + \mu_i^X, \epsilon_i \in \mathbb{R}^d;$$

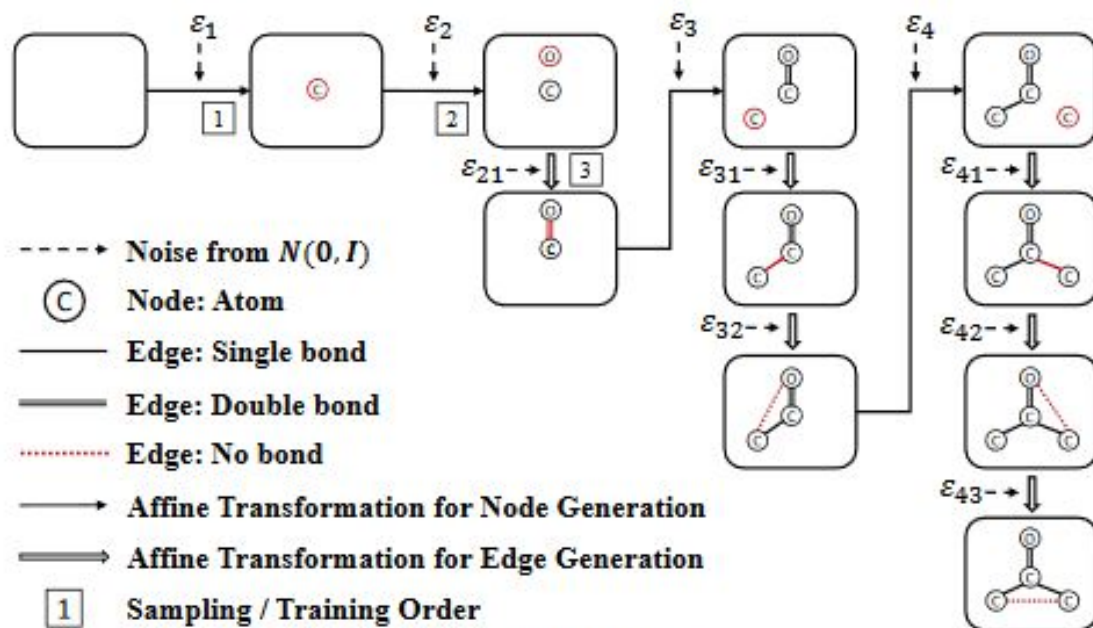
$$z_{ij}^A = \epsilon_{ij} \odot \alpha_{ij}^A + \mu_{ij}^A, j \in \{1, 2, \dots, i-1\}, \epsilon_{ij} \in \mathbb{R}^{b+1}$$

- Discretize the real value graph  $z = (z^A, z^X)$  to  $G = (A, X)$  by using argmax

$$X_i = v_{\text{argmax}(z_i^X)}^d$$

$$A_{ij} = v_{\text{argmax}(z_{ij}^A)}^{b+1}$$

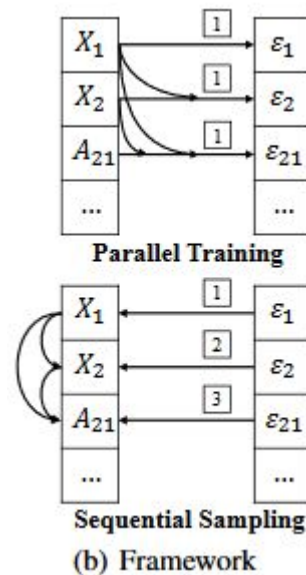
## Methods: GraphAF



(a) Sampling Phases

## Methods: Efficient Training

- Jacobian of the inverse mapping of normalizing flow is lower triangular
  - Determinant of lower triangular jacobian can be efficiently computed
- Masking nodes and edges in a graph can be used to train R-GCN in parallel
  - Satisfy the autoregressive property by forcing every masked graph to be a subgraph of previous graph
    - Use  $G_i$  to predict  $\epsilon_i$  and  $G_{i-1}$  to predict  $\epsilon_{i-1}$  and  $G_{i-1} \subset G_i$







# Quiz

SC: Why GraphAF training efficient?

- A. The Jacobian of the inverse of normalizing flow is sparse
- B. GraphAF utilizes transformers for graph generation which is more parallelizable than RNNs
- C. **Iterations of the graph generation can be trained independently because masking satisfies the autoregressive property**

# Methods: Chemistry Recap

- Atoms need to have a specific number of electrons in their outermost orbital to be stable
  - Number of electrons it can give up or accept is known as the valence number
  - For nonmetals and some metals, valence number can be determined by the group number
  - For most metals, valence number can change based on the type of bond formed

**+1 VALENES (OXIDATION NUMBERS) OF A-GROUP ELEMENTS 0**

1A	2A	3A	4A	5A	6A	7A	8A
1 H 1.01	4 Be 9.012	5 B 10.81	6 C 12.01	7 N 14	8 O 16	9 F 19	10 Ne 20.18
11 Na 22.99	12 Mg 24.31	13 Al 26.98	14 Si 28.1	15 P 30.97	16 S 32.06	17 Cl 35.45	18 Ar 39.95
19 K 39.10	20 Ca 40.08	21 Sc 44.96	22 Ti 47.9	23 V 50.94	24 Cr 51.99	25 Mn 54.93	26 Fe 55.85
27 Co 58.93	28 Ni 58.71	29 Cu 63.54	30 Zn 65.37	31 Ga 69.72	32 Ge 72.59	33 As 74.92	34 Se 78.96
37 Rb 85.47	38 Sr 87.62	39 Y 88.91	40 Zr 91.22	41 Nb 92.9	42 Mo 95.94	43 Tc 99	44 Ru 101.07
47 Ag 107.9	48 Cd 112.4	49 In 114.8	50 Sn 118.7	51 Sb 121.7	52 Te 127.6	53 I 126.9	54 Xe 131.3
55 Cs 132.9	56 Ba 137.3	57 La 138.9	58 Ce 140.1	59 Pr 140.9	60 Nd 144.2	61 Pm 144.9	62 Sm 150.4
69 Tm 168.9	70 Yb 173.0	71 Lu 175.0	72 Hf 178.5	73 Ta 180.9	74 W 183.85	75 Re 186.2	76 Os 190.2
77 Ir 192.2	78 Pt 195.1	79 Au 196.9	80 Hg 200.6	81 Tl 204.4	82 Pb 207.2	83 Bi 209	84 Po 210
87 Fr 223	88 Ra 226	89 Ac 227	90 Th 232	91 Pa 231	92 U 238	93 Np 237	94 Pu 244

Can't tell by looking at the table  
Transition Metals

lose 1e- lose 2e- lose 3e- lose or gain 4e- gain 3e- gain 2e- gain 1e- nada!

Metals  
Nonmetals  
Metalloids



## Methods: Validity Constraints

- Valence number can be used as a constraint to determine if a molecule is valid

$$\sum_j |A_{ij}| \leq \text{Valency}(X_i) \text{ and } \sum_i |A_{ij}| \leq \text{Valency}(X_j)$$

- If newly added edge violates the valency constraint, reject the edge



# Quiz

SC: How does GraphAF handle valency constraints in the generation process?

- A. Use the constraint as a term in the loss function and optimize using gradient descent
- B. **At every step in sequential generation, check if the newly generated molecule obeys the valency constraint**
- C. After the molecule has been generated, use a pruning algorithm to remove excess edges



# Methods: Property Optimization

- Use reinforcement learning to choose nodes and edges that meet a given criteria
- State
  - The current subgraph
  - Only allow state transitions that satisfy the valency constraint
- Policy
  - The autoregressive model defined earlier
  - Given the current state, the subgraph, choose the best action, which node and edges to add



# Methods: Property Optimization

- Reward Function
  - Utilizes intermediate rewards, calculated when an action is taken, and final rewards, calculated when the generation procedure is complete
    - Intermediate reward: negative reward if immediate action violates valency rule
    - Final reward: Desired property score, logP, and drug likeness, QED
      - Distributed to intermediate steps with a discount factor

$$L(\theta) = -E_{G \sim p_\theta} \left\{ E_i \left[ \min(r_i(\theta)V(G_i, X_i), \text{clip}(r_i(\theta), 1 - \epsilon, 1 + \epsilon)V(G_i, X_i)) \right. \right. \\ \left. \left. + E_j [\min(r_{ij}(\theta)V(G_{ij}, A_{ij}), \text{clip}(r_{ij}(\theta), 1 - \epsilon, 1 + \epsilon)V(G_{ij}, A_{ij}))] \right] \right\}$$

- where  $V(\text{state}, \text{action})$  is the advantage function,  $r_i(\theta) = \frac{p_\theta(X_i|G_i)}{p_{\theta_{old}}(X_i|G_i)}$ ,  $r_{ij}(\theta) = \frac{p_\theta(A_{ij}|G_{ij})}{p_{\theta_{old}}(A_{ij}|G_{ij})}$  and  $G_{ij} = G_i \cup X_i \cup A_{i,1:j-1}$



# Experiments

## Evaluation Tasks

1. Density Modeling and Generation: learn the data distribution and generate realistic and diverse molecules
2. Property Optimization: generating novel molecules with optimized chemical properties
3. Constrained Property Optimization: modifying the given molecule to improve desired properties while satisfying a similarity constraint

**Data** ZINC250k molecular dataset: 250,000 drug-like molecules

**Baselines** JT-VAE (Jin et al., 2018), GCPN, MolecularRNN (Popova et al., 2019), GraphNVP (Madhawa et al., 2019)



# Metrics

**Validity** is the percentage of valid molecules among all the generated graphs.

**Uniqueness** is the percentage of unique molecules among all the generated molecules.

**Novelty** is the percentage of generated molecules not appearing in training set.

**Reconstruction** is the percentage of the molecules that can be reconstructed from latent vectors.

The above metrics are calculated from 10,000 randomly generated molecules.





## Density Modeling and Generation

Table 2: Comparison of different models on density modeling and generation. *Reconstruction* is only evaluated on latent variable models. *Validity w/o check* is only evaluated on models with valency constraints. Result with  $\dagger$  is obtained by running GCPN’s open-source code. Results with  $\ddagger$  are taken from [Popova et al. \(2019\)](#).

Method	Validity	Validity w/o check	Uniqueness	Novelty	Reconstruction
JT-VAE	100%	—	100% $\ddagger$	100% $\ddagger$	76.7%
GCPN	100%	20% $\dagger$	99.97% $\ddagger$	100% $\ddagger$	—
MRNN	100%	65%	99.89%	100%	—
GraphNVP	42.60%	—	94.80%	100%	100%
GraphAF	100%	68%	99.10%	100%	100%

- GraphAF (as flow based model) holds perfect reconstruction ability compared with VAE approaches.
- GraphAF achieves a 100% validity rate due to the valency check during sequential generation (GraphNVP only achieves 42.60% due to its one-shot sampling process)
- JT-VAE and GCPN take around 24 and 8 hours, respectively, while GraphAF only takes 4 hours




# Property Optimization

Table 5: Comparison of the top 3 property scores of generated molecules.


Method	Penalized logP				QED			
	1st	2nd	3rd	Validity	1st	2nd	3rd	Validity
ZINC (Dataset)	4.52	4.30	4.23	100.0%	0.948	0.948	0.948	100.0%
JT-VAE (Jin et al., 2018)	5.30	4.93	4.49	100.0%	0.925	0.911	0.910	100.0%
GCPN (You et al., 2018a)	7.98	7.85	7.80	100.0%	<b>0.948</b>	0.947	0.946	100.0%
MRNN <sup>1</sup> (Popova et al., 2019)	8.63	6.08	4.73	100.0%	0.844	0.796	0.736	100.0%
GraphAF	<b>12.23</b>	<b>11.29</b>	<b>11.05</b>	100.0%	<b>0.948</b>	<b>0.948</b>	<b>0.947</b>	100.0%

- logP score penalized by ring size and synthetic accessibility.
- QED measures the druglikeness of the molecules.
- Pretrain the GraphAF network for 300 epochs for likelihood modeling, and apply the RL process to fine-tune the network


# Property Optimization



(a) Penalized logP optimization




0.948



0.947



0.948



0.946

(b) QED optimization



# Property Optimization

GraphAF v.s. GCPN

- RL process is similar
- GCPN uses GAN model which is known to suffer from the mode collapse problem
- GraphAF uses flow which is flexible at modeling complex distribution and generating diverse data, encouraging exploration during RL



# Constrained Property Optimization

Table 6: Comparison of results on constrained property optimization.

$\delta$	JT-VAE			GCPN			GraphAF		
	Improvement	Similarity	Success	Improvement	Similarity	Success	Improvement	Similarity	Success
0.0	$1.91 \pm 2.04$	$0.28 \pm 0.15$	97.5%	$4.20 \pm 1.28$	$0.32 \pm 0.12$	100%	<b><math>13.13 \pm 6.89</math></b>	$0.29 \pm 0.15$	100%
0.2	$1.68 \pm 1.85$	$0.33 \pm 0.13$	97.1%	$4.12 \pm 1.19$	$0.34 \pm 0.11$	100%	<b><math>11.90 \pm 6.86</math></b>	$0.33 \pm 0.12$	100%
0.4	$0.84 \pm 1.45$	$0.51 \pm 0.10$	83.6%	$2.49 \pm 1.30$	$0.47 \pm 0.08$	100%	<b><math>8.21 \pm 6.51</math></b>	$0.49 \pm 0.09$	99.88%
0.6	$0.21 \pm 0.71$	$0.69 \pm 0.06$	46.4%	$0.79 \pm 0.63$	$0.68 \pm 0.08$	100%	<b><math>4.98 \pm 6.49</math></b>	$0.66 \pm 0.05$	96.88%

- Optimize penalized logP for 800 molecules in ZINC250k with the lowest scores
- Set the initial states as sub-graphs randomly sampled from 800 molecules to be optimized during generation



## Conclusion

- GraphAF is the first flow-based autoregressive model for generating realistic and diverse molecular graphs.
- GraphAF generates novel and 100% valid molecules in empirical experiments, and its training process is fast.
- The generative process is fine-tuned with reinforcement learning to optimize the properties of generated molecules,